

LEARNSLICE GUIDE

# AI for product managers

97% of product managers say AI made them faster. Only 64% say it made their product better.

---

Inside: why your Copilot probably cannot see your work, the five places AI earns its keep in a product week, and the prompts worth saving.



# Contents

---

If you only have twenty minutes

---

Half the profession is stuck in the same place

---

## Part 0: six things nobody told you

1. Your best prompt might have been luck
  2. Your Copilot probably cannot see your work
  3. The four-part prompt Microsoft wrote and nobody reads
  4. The one key most people never find
  5. Where your prompts actually go
  6. It will lie to you in your own house style
- 

## Part 1: four habits that separate method from luck

- Rule 1. Give it your own material
  - Rule 2. Know what must never go in
  - Rule 3. Check it before it leaves you
  - Rule 4. Keep what works
- 

## Part 2: five places it earns its keep

- Chapter 1. Why Wednesday never looks like Monday's plan
  - Chapter 2. You know the product better than the people using it
  - Chapter 3. Everything is important, and that is the problem
  - Chapter 4. Most of your week is writing nobody acts on
  - Chapter 5. The metric moved and the meeting is tomorrow
- 

## Part 3: the prompt pack

---

## Part 4: staying current on thirty minutes a week

- Thirty minutes a week
  - Where to look, and how often
  - The weekly
- 

If you lead a product team

---

Sources

---

What we do


---

That gap is what this guide is about. Not how to write a clever prompt, but how to get from working quicker to shipping something a stakeholder would notice.

## Faster, but not better

Product managers reporting improvement from AI

Personal productivity



Product outcomes



**33 percentage points between working quicker and shipping better**

Every number here comes from one source: the Product Focus 2026 Survey of the Product Management Profession, 677 respondents across 40 countries, collected between October 2025 and January 2026. Most respondents were in Europe, 83%, with 8% in the United States. Anything about how Microsoft Copilot behaves comes from Microsoft's own documentation and was checked in September 2026. We are not affiliated with either.

---

## If you only have twenty minutes

If you have barely used AI at work, skip this and read the next section instead. Otherwise, read four things and stop.

- 1. Your Copilot probably cannot see your work,** Part 0, item 2. Most people who think Copilot is useless are on a tier that cannot read their work.
- 2. Give it your own material.** Rule 1. This is the difference between a generic answer and a useful one, and it is not about phrasing.
- 3. Check it before it leaves you.** Rule 3. Numbers, names, and anything stated as a fact about a customer.
- 4. One chapter,** whichever of the five describes your worst week.

Everything else is depth you can come back for.

---

## Half the profession is stuck in the same place

Half of the product managers in this survey who do not use AI say the reason is that they are unsure how. More than either distrust it, at 33%, or cite legal and security concerns, at 17%. Those three figures describe the minority who stay away, and the report does not state how many people that is, so read them as a ranking of reasons rather than as a precise count. The ranking is the useful part: not knowing how comes first. If that is you, you are in the largest group, not behind it.

What the tool is, in two sentences: a system that predicts likely text from the text you give it, trained on an enormous amount of writing. That is why it is fluent, why it is useful on anything language-shaped, and why it will state something false with complete confidence.

**Start here, once, this week.** Take a document you already wrote and know well. Ask it to list the questions a sceptical reader would ask about it. You can judge that answer immediately, which is the point of starting on work you already understand. Then read the rest of this.

---

---

## Part 0: six things nobody told you

Six short sections before the guide proper. Each one takes a minute to read and heads off a specific frustration that otherwise looks like your own fault.

### 1. Your best prompt might have been luck

Microsoft says it plainly in its own guidance: using the same prompt multiple times can result in different responses. This is how the technology works, not a fault.

The practical consequence is the one people miss. **One good answer does not prove the prompt was good.** It might have been luck. A prompt is only worth keeping once it has produced a usable answer more than once, on different inputs. That is why Rule 4 exists.

### 2. Your Copilot probably cannot see your work

There are three tiers and they behave completely differently.

---

## What each Copilot can reach

Same name, three different products

|                                                 | WEB         | IN APPS     | YOUR FILES  |
|-------------------------------------------------|-------------|-------------|-------------|
| Copilot Chat<br>(Basic)                         | <div></div> | <div></div> | <div></div> |
| Microsoft 365 Copilot<br>(Basic)                | <div></div> | <div></div> | <div></div> |
| Microsoft 365 Copilot<br>(Premium), paid add-on | <div></div> | <div></div> | <div></div> |

An empty cell means it sees your work only if you paste or upload it.

The same thing in detail, including what each tier can reach and how:

| TIER                                             | CAN IT SEE YOUR FILES AND MAIL                                                                            |
|--------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| Copilot Chat (Basic)                             | No. Web data only. It sees your work only if you paste it, upload it, or have it open in Teams or Outlook |
| Microsoft 365 Copilot (Basic)                    | Not in chat. But Copilot works inside Word, Excel, PowerPoint and OneNote                                 |
| Microsoft 365 Copilot (Premium), the paid add-on | Yes, automatically, through Microsoft Graph, and only for files you already have permission to open       |

If you have asked Copilot "what did we decide in the roadmap review" and got nothing useful, you are almost certainly not on the premium tier. That is a licensing fact, not a skill problem.

**The model picker.** Copilot offers Auto, Quick response and Think deeper. Auto chooses for you. Switch to Think deeper when the task is a judgment rather than a lookup, and expect it to take longer on purpose.

**Grounding, which is the word for all of this.** Two sources feed any answer: the public web, and your organisation's own content. Only the premium tier reaches the second one by itself. Every other tier sees only what you hand it.

**Which tool, not which winner.** In this profession 66% of AI users use ChatGPT regularly, with Copilot second, then Gemini, Claude and Perplexity. 22% use three or more regularly, and 52% of ChatGPT users

also use Copilot. Nobody getting value out of this is loyal to one tool. They pick per task: the one that can see your work for anything about your work, a general chat tool for thinking out loud.

**Before you look for a new tool.** Jira is recommended by 58% of this profession and Confluence by 46%, so for most readers the product data already sits in tools that are adding their own AI on top of it. Check what you already pay for before buying anything new.

### 3. The four-part prompt Microsoft wrote and nobody reads

This is Microsoft's own framework, and it is a good one.

**Goal.** What you want. **Context.** The situation, who it is for, what has already been decided. **Expectations.** Format, length, tone, what not to do. **Source.** Where it should get the material.

A thin prompt: *Summarise this feedback.*

The same prompt with four parts: *Summarise the customer feedback in the file below. Context: I am preparing a prioritisation session for a B2B product and I need themes, not individual complaints. Expectations: at most six themes, each with a count and one verbatim quote, no recommendations. Source: only the attached file.*

The second one is not more clever. It is more specific about what you already know.

### 4. The one key most people never find

In Copilot, type  and start typing the name of a file, person, meeting or email. You can attach a single file or a whole folder. Inside a SharePoint site you can reference up to ten files or pages. Checked September 2026.

Most people paste content they could have referenced, and lose the formatting and the source link doing it.

### 5. Where your prompts actually go

When you are signed in with your work account, prompts and responses are covered by enterprise data protection, and Microsoft states they are not used to train the foundation models. Access is scoped by your existing permissions, so Copilot cannot show you a document you could not already open yourself. For users in the European Union there are additional European Union Data Boundary safeguards.

**The caveat that matters.** All of that describes a work account signed in with your company identity. A personal account is a different product with different terms. If you are not certain which you are signed into, check before you paste anything that matters.

This is practical guidance, not legal advice. Your company's own rules sit on top of it.

## 6. It will lie to you in your own house style

It will state things that are not true, confidently and in your own house style. Three things need checking every single time:

- **Numbers.** Check every figure you did not give it yourself. It may have retrieved a real one, and it may have produced something that merely looks like one.
- **Names.** People, products, companies, competitors.
- **Anything stated as a fact about a customer.** This is the dangerous one, because it is the hardest to spot and the most expensive to act on.


Microsoft's own guidance says to review and verify responses. Take the vendor at its word.

**What this looks like in practice.** You paste thirty support tickets and ask for the themes. Back comes: "Customers consistently request single sign-on." You go looking, and one ticket mentioned it once, as an aside. Nothing was invented exactly. A single signal was promoted to a pattern, in confident language, in a sentence you could reasonably paste into a roadmap review. That is the failure mode to expect. Not a made-up fact you would catch, but a real thing overstated, reading perfectly.

## One mention is not a pattern

## The failure mode to expect when you ask for themes

## 30 SUPPORT TICKETS PASTED IN



## WHAT CAME BACK


"Customers consistently request single sign-on."

One ticket in thirty, reported as consistent.

**The decision aid.** Before you use an output, ask one question: *could I tell if this were wrong?* If yes, use it and check it. If no, you are not the right person to accept that output yet, and neither is the tool. Get it from someone who would know, or go and find out. It is one question and it catches the outputs that are worth worrying about.

## Before you use the output

## One question, asked every time



## Part 1: four habits that separate method from luck

Part 0 was how the tool behaves. These four are how you behave, and they are the whole difference between an answer that worked once and a result you can repeat.

## Rule 1. Give it your own material

The difference between a generic answer and a useful one is almost always what you put in, not how you phrase the question.

Paste the actual ticket, not your summary of it. The actual transcript, not your memory of the call. The actual numbers, not "roughly a third". A model working from your description of a thing is working from a copy of a copy.

## Rule 2. Know what must never go in

Among product managers who avoid AI, 17% name legal or security reasons. They are not being irrational, they are being unspecific. Turn it into three questions you can answer in a second:

1. **Which account am I in?** Work identity, or personal. See Part 0, item 5.
2. **Would I put this in an email to a supplier?** If not, it does not go in a chat window either.
3. **Is it someone else's personal data?** Customer names, recordings, support tickets with identities in them. Anonymise or do not paste.

Your company's policy overrides all three. If it has one, read it once, properly. If it does not, ask whoever would own it.

## Rule 3. Check it before it leaves you

85% of this profession already validate AI output using their own expertise, so this rule confirms a habit rather than introducing one. What most people do not have is a consistent list. Use the three from Part 0, item 6: numbers, names and customer claims.

The moment that matters is **before it leaves you**, not before it ships. Once a plausible wrong sentence is inside a document three other people have edited, it stops being an AI problem and becomes a fact your organisation believes.

## Rule 4. Keep what works

The habit with the most compounding value, and the easiest one to skip.

When an answer comes out genuinely good, do not just use it. Save the prompt that produced it, with the parts that change marked in square brackets, in whatever you already use: a note, a Confluence page, a

pinned message. Next time the same job takes a minute instead of twenty.

---

## Part 2: five places it earns its keep

Rules are easy to agree with and easy to forget. Here they are in the five places a product week actually bends, ordered by how much the profession says each one hurts, so firefighting comes first.

### Chapter 1. Why Wednesday never looks like Monday's plan

**The need.** You planned the week on Monday and by Wednesday you are working on none of it.

**The number.** 60% of product managers say unplanned work frequently disrupts their schedule. Firefighting heads the survey's list of the big issues the profession faces, ahead of lack of resource and weak or missing company strategy.

#### Where it genuinely helps.

- Turning fifty inbound requests into six themes, so the argument is about categories rather than tickets.
- Drafting the reply that declines a request and still explains the reasoning. This is the message most people avoid writing, and therefore never send.
- Capturing a corridor decision into three lines that go into the ticket before it evaporates.

**What it cannot do.** It cannot reduce the number of interruptions. Firefighting tracks the environment, not the person: in this survey automotive reported 76% while SaaS reported 57%. If you are interrupted constantly, the fix is a conversation with your manager about routing, and no tool substitutes for it. What you can change is how much of your week each interruption costs after it arrives.

**One prompt.** Works on any tier, because you paste the content in.

Goal: group the requests below into at most six themes.

Context: these are [N] requests that came to me this week as a product manager for [product].

Expectations: for each theme give the theme in one line, how many requests it covers, and the single question I would need answered to decide what to do about it. Mark any theme that rests on only one request. Do not suggest solutions. Quote the requests, do not paraphrase them.

Source: only the text pasted below.

[paste the requests]

**Your first step, under fifteen minutes.** Take this week's inbound, in whatever form it is in, and run that prompt once. If the six themes come out wrong, that tells you something too. Usually it means the requests are not really requests, they are escalations.

## Chapter 2. You know the product better than the people using it

**The need.** Every call you make is a guess about someone you have not spoken to in weeks.

**The number.** 71% of product managers say they do not spend enough time with customers or understanding the market. It is worse the more senior you get: 75% among heads and directors, against 59% among junior product managers.

### It gets worse as you get senior

Say they do not spend enough time with customers



**The people who could fix it are the furthest from it.**

### Where it genuinely helps.

- Turning interview recordings or notes into themes with the original quotes still attached, so you can check the theme against the words.
- Preparing sharper questions from what you already know, so the half hour you do get is not spent on things you could have read.
- Finding the contradictions between what two customers told you.

**What it cannot do, and this one is serious.** It cannot talk to a customer for you, and it must never be used to invent one. Asking a model to "act as our typical user" and answer your questions produces fluent, confident, plausible fiction. It is the single most damaging misuse available to this profession, because the output looks exactly like research and carries none of the risk that real research carries: the risk of being told you are wrong.

### One prompt.

```
Goal: extract the themes from the customer interview transcript
below.

Context: [product], [segment]. I am looking for problems the
customer
described, not features they requested.

Expectations: at most five themes. For each: the theme in one
line, and
up to two verbatim quotes that support it, nothing paraphrased.
If only
one quote supports a theme, give one and say so. Do not
describe
anything mentioned once as a pattern. Then a separate list of
anything
the customer said that contradicts another part of the same
interview.

Source: only the transcript below.

[paste the transcript]
```

**Your first step.** Take the most recent customer conversation you have any record of, even rough notes, and run it. The contradictions list is usually the part that earns the fifteen minutes.

## Chapter 3. Everything is important, and that is the problem

**The need.** You cannot say no, because there is nothing concrete to say no against.

**The numbers.** 33% of product managers report a weak or missing company strategy. 34% have no clear primary metric they are accountable for. Those two travel together: product managers who report a weak company strategy are 15 percentage points more likely to lack a primary metric, 44% against 29%. It is difficult to prioritise if nobody has defined what winning looks like.

### Where it genuinely helps.

- Turning a vague strategy statement into a written test you can actually apply, with the questions it would have to answer to decide a real case.
- Making the trade-off in a decision explicit, so the thing you are giving up is on the page rather than discovered later.
- Arguing the other side of your own case before somebody else does it in the meeting.

**What it cannot do.** It cannot supply the strategy. It will happily generate a plausible one, and that is the trap: a fluent strategy nobody agreed to is worse than an admitted absence, because it ends arguments that needed to happen. Use it to sharpen a strategy that exists, or to make the absence of one visible. Never to fill the hole quietly.

### One prompt.

```
Goal: turn the strategy statement below into a prioritisation
test.
Context: I need to decide between competing requests and I want
a test
I can apply consistently and defend in a meeting.
Expectations: five questions, each answerable yes or no about a
specific
request. Then tell me which parts of the strategy statement are
too vague
to generate a question from, and say why.
Source: only the statement below.

[paste your company or product strategy as written]
```

**Your first step.** Run it on your actual strategy, then apply the five questions to the two requests currently competing for the same sprint. The list of what was too vague to test is the more valuable half of the output, and it is the thing to bring to your manager.

## Chapter 4. Most of your week is writing nobody acts on

**The need.** Most of the job is writing, and most of the writing is read by someone who has thirty seconds.

**The number.** Asked which activity they spend the most time on, 56% of product managers name inbound activities, against 25% naming strategic activities and 19% outbound. A great deal of that inbound work is documents, updates, tickets and mail.

### Where the week actually goes

Share of product managers naming this as where most of their time goes



Most of the job is inbound, and most of inbound is writing.

### Where it genuinely helps.

- The second draft, not the first. Write the bad version yourself in five minutes, then hand it over to be tightened. The first draft carries your intent and that is the part a model cannot guess.
- Compressing a long document down to the decision it is asking for.
- Rewriting one update for three audiences, because the engineering version and the executive version of the same news are genuinely different documents.

**What it cannot do.** It cannot know what you meant. A first draft generated from an empty prompt costs more to fix than it would have cost to write, because you end up editing someone else's argument instead of making your own.

### One prompt.

Goal: rewrite the update below for three audiences.

Context: same news, three readers: the engineering team, my head of product, and a customer-facing colleague who needs to answer questions about it.

Expectations: three versions, each under 120 words. Each opens with what the reader must do or decide. Do not add any information that is not in my draft. Flag anything in my draft that is too vague to translate.

Source: only my draft below.

[paste your draft]

**Your first step.** Take the last update you sent that got no reply, and run it. What comes back as "too vague to translate" is usually the reason it got no reply.

## Chapter 5. The metric moved and the meeting is tomorrow

**The need.** You have the number. You do not know which of six explanations is the right one, and you are presenting it tomorrow.

**The number.** When product managers were asked which skills will matter most over the next two years, data analysis and literacy was named second among hard skills, behind only AI proficiency itself.

### Where it genuinely helps.

- Drafting the query or the formula, which is a language problem more than a maths one.
- Listing the plausible explanations for a movement, including the boring ones like a reporting change or a seasonal effect, so you walk in having ruled things out.
- Telling you what is missing from an analysis before you present it, which is a much better use than asking it to do the analysis.

**What it cannot do.** It cannot be trusted on arithmetic, and it cannot be trusted on any number it was not given. Every figure that leaves you is yours, not its. If you did not check it, you did not say it, you repeated it.

## One prompt.

```
Goal: list the plausible explanations for the change below.
Context: [metric] moved from [x] to [y] between [date] and
[date] for
[product]. Known changes in that window: [releases, campaigns,
pricing,
seasonality, reporting changes].
Expectations: rank the explanations by how easily each could be
checked,
cheapest first, and for each name the exact check. Include
mundane
explanations such as instrumentation or reporting changes. Do
not
calculate anything I have not given you.
Source: only what I have written above.
```

**Your first step.** Take the last metric movement you had to explain and run it. If the cheapest check on the list is one nobody ran, that is your meeting.

---

## Part 3: the prompt pack

Every prompt below uses the four parts from Part 0, item 3. Square brackets mark the parts you change. Each says what it needs, and "any tier" means it works even on Copilot Chat (Basic), because you supply the content yourself.

### COPILOT CHAT, ANY TIER

---

Theme a pile of requests: see Chapter 1.

## Argue the other side.

```
Goal: make the strongest case against the decision below.
Context: [decision], [what it costs], [who disagrees].
Expectations: five objections, strongest first, each with the
evidence someone would need to defeat it.
Source: only what I have written.
```

## Prepare for a difficult conversation.

Goal: list the questions [role] will ask about [decision].  
Expectations: the five hardest. Give the honest answer to each where I have given you enough to answer, and say plainly "you do not have this" where I have not.  
Source: only what I have written.

## OUTLOOK, NEEDS MICROSOFT 365 COPILOT

---

Draft the decline: see Chapter 1.

## Find the decision buried in a long thread.

Goal: tell me what this thread is actually asking me to decide.  
Expectations: the decision in one line, who is waiting, what is blocking it, and anything decided earlier in the thread that people seem to have forgotten.  
Source: this thread only.

## TEAMS, NEEDS MICROSOFT 365 COPILOT

---

Copilot in Teams reaches back over the last 30 days of meetings, which is easy to miss.

## Turn a recap into owners and dates.

Goal: extract the commitments made in this meeting.  
Expectations: a table of commitment, owner, date. Anything without a named owner goes in a separate list titled unassigned. Do not invent owners or dates.  
Source: this meeting only.

**Get what was left unresolved**, which the standard recap tends to leave out.

Goal: list what this meeting did not settle.  
Expectations: open questions and disagreements only, no summary of  
what was agreed. Name who raised each one.  
Source: this meeting only.

#### WORD, NEEDS MICROSOFT 365 COPILOT

---

Rewrite for three audiences: see Chapter 4.

**Compress a document to the decision it asks for.**

Goal: tell me what this document is asking the reader to decide.  
Expectations: the decision in one line, the three facts that bear  
on it, and anything the document assumes without stating.  
Source: this document only.

#### EXCEL, NEEDS MICROSOFT 365 COPILOT

---

Explain a metric movement: see Chapter 5.

**Find what the analysis cannot answer.**

Goal: tell me what is missing from this analysis.  
Expectations: list the questions a sceptical reader would ask that  
this data cannot answer. Do not attempt to answer them, and do not  
calculate anything new.  
Source: the selected data only.

#### POWERPOINT, NEEDS MICROSOFT 365 COPILOT

---

**Turn a decision document into slides.**

Goal: turn the decision document below into five slides.  
Expectations: one slide per argument, the recommendation on slide one and not slide five, no slide with more than three bullets.  
Source: only the document below.

## Part 4: staying current on thirty minutes a week

93% of this profession say they want to learn more about AI tools. The obstacle is not appetite. 48% say a lack of budget prevents them getting training and 23% say there is no management support for it, and only 32% rate the development opportunities at their organisation as better than average.

So here is a routine that costs nothing and needs nobody's approval.

If your company has deployed nothing at all, it still works. Use a free tier with a personal account, never put company material into it, see Rule 2, and practise on public material or your own writing. The habits transfer. The grounding does not.

### Thirty minutes a week

#### Thirty minutes a week

No budget, no approval, no course



##### 10 min

What changed.  
One source,  
rotating.

##### 15 min

A real task from your week,  
done again with the tool.  
You already know what  
good looks like.

##### 5 min

Write it  
down.

The middle block is the one that teaches. The other two make it stick.

**Ten minutes, what changed.** Check one source from the list below. Not all of them, one, rotating.

**Fifteen minutes, on your own real work.** Take a task you did this week and do it again with the tool. Not an exercise, the actual task, where you

already know what good looks like. This is the only way to tell whether the output is any good, and it is why learning on your own work beats any course.

**Five minutes, write it down.** One line in your saved prompts: what worked, what did not. This is Rule 4, and it is what turns thirty minutes into something cumulative.

## Where to look, and how often

| QUESTION                             | SOURCE                                                                                                      | HOW OFTEN            |
|--------------------------------------|-------------------------------------------------------------------------------------------------------------|----------------------|
| What changed in Copilot              | Microsoft 365 Copilot release notes on Microsoft Learn, and the Microsoft 365 roadmap                       | Monthly              |
| What changed in the models           | The providers' own release notes and model cards                                                            | When something ships |
| What changed in my profession        | The Product Focus annual survey, the source of this guide                                                   | Yearly               |
| What changed in the rules, in the EU | The European Commission's AI Act pages. The AI literacy duty in Article 4 has applied since 2 February 2025 | Quarterly            |

One policy worth adopting: **do not treat vendor blogs, news aggregators or social threads as primary sources.** Every number in this guide had to survive being traced back to the organisation that published it. Most of what circulates does not survive that.

## The weekly

If you would rather not do the first ten minutes yourself, we do it and send one email a week. One change that matters and what to do about it, one prompt of the week, one thing worth reading and why, and one line on what we got wrong the week before. Five minutes. Unsubscribe whenever.

Sign up at [learnslice.com/ai-guides](https://learnslice.com/ai-guides). If this was forwarded to you, that is also where the other guide is.

---

## If you lead a product team

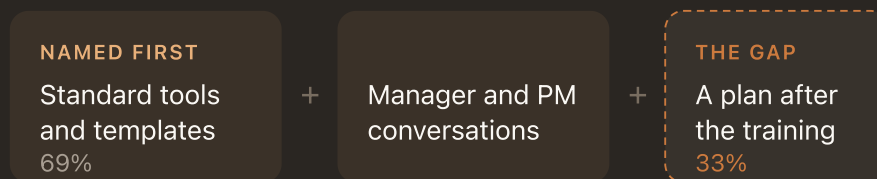
41% of respondents manage a team of product people, so this section is for a large minority of readers.

Product leaders in the survey were asked what actually embeds new product skills and behaviour in the workplace. Standard tools and templates came first, named by 69%, ahead of structured meetings between line managers and product managers, and ahead of standardised training for teams. Developing a prioritised plan after training was named by 33%.

The survey's own conclusion is the useful part: the best results need all three elements together, training and tools, structured manager conversations, and an activation plan after the training. Most organisations do the first and skip the third.

### Three things, or it does not stick

What leaders say embeds a new product skill



---

Most organisations do the first and skip the third.

**The third is the one this guide hands you for free.**

Three practical moves, in that order:

- 1. Make the prompt pack the team's shared templates.** Not a document nobody opens, a place where anyone can add one that worked. This is the 69% finding applied.
- 2. Put one question in the one to one.** "What did you try with AI this month and what came out wrong?" The second half is what makes it safe to answer honestly.
- 3. Run a four week activation plan after anyone learns anything.** Week one, everyone picks a recurring task. Week two, they do it with the tool and keep the prompt. Week three, prompts go into the shared place. Week four, one person shows what did not work. This is the element only a third of leaders have.

If you want the four weeks run for the team rather than assembled by you, that is what the last page is about.

One last thing. Chapter 2 showed that the customer-time deficit is worst at your level, not your team's. Whatever the three moves above free up, that is where it should go.

---

---

## Sources

- **Product Focus, 2026 Survey of the Product Management Profession.** 677 respondents across 40 countries, collected October 2025 to January 2026. 83% Europe, 8% United States, 9% elsewhere. Every figure about the profession in this guide comes from this report. We are not affiliated with Product Focus and reproduce none of its charts.
- **Microsoft**, product documentation on Microsoft Learn and Microsoft Support, for Copilot licence tiers, grounding, data protection, file referencing, prompt structure and per-application features. Checked September 2026. Copilot changes often, which is what the weekly is for.
- **European Commission**, for the AI Act and the Article 4 AI literacy duty applicable since 2 February 2025.

No figure in this guide is estimated, modelled or taken from a vendor's marketing material. Where we give advice rather than a finding, it is written as advice.

## What we do

Nothing before this page was an advertisement. This page is, and it is the only one.

**When the limit is the tool, not you.** Every chapter above ends with something the tool cannot do, and most of those limits are the same limit: the model cannot read your own material. It cannot search your tickets, your documents or your decisions, so it answers from the public web and from whatever you paste. We build the layer that fixes that, grounded on a company's own content, hosted where the company needs it hosted. If that is the wall you keep hitting, it is worth a conversation.

**AI Mentor, for a team rather than a person.** The routine and the activation plan in this guide work for one person with thirty minutes a week. Running the same thing across a team, with content matched to your sector and a record of who actually learned what, is what AI Mentor does. It is the same discipline as this guide, delivered as a product.

If you picked one of those two on the download form, we have your answer and will come back to you on it. If you picked neither and changed your mind, or this was forwarded to you, [learnslice.com/ai-guides](https://learnslice.com/ai-guides) is where to say so.